

## Final-Submission – Logic Explanation

### NOTES:

1. I have submitted below artifacts in final submission for capstone project:
  1. FraudAnalysis.zip
  2. LogicFinal.pdf
  3. Output.pdf
  4. Output\_without\_details\_of\_records.pdf
  5. Output.txt
  6. card\_transactions.csv

I have deliberately provided **Output\_without\_details\_of\_records.pdf**. I have displayed all 6075 records received from Kafka topic in the Output.pdf file. Just in case, evaluator directly wants to see logs for processing of transactions then can open Output\_without\_details\_of\_records.pdf or evaluator will have to scroll down up to page number 290 to start with the logs for processing of transactions.

I have deliberately provided **Output.txt** as well. This is the original file which captured the output of spark program. Since line length of a text file is much higher than of a word or pdf file, if you open this Output.txt file in Notepad++ then you will see output with more clarity.

**Please zoom to 180% or 200% to see screenshots with better clarity.**

2. As per submission guidelines for final submission point 8, we are supposed to submit screenshots of console output. Since program was running very fast, I was not able to take screenshots. I was able to redirect console output to a file. It has been confirmed by TA in this discussion forum link that we can submit that output file and it will suffice: <https://learn.upgrad.com/v/course/119/question/129993>
3. This link mentioned in resources section of capstone project, was used for Kafka configuration: <https://learn.upgrad.com/v/course/119/question/123157>
4. This link mentioned in resources section of capstone project, was used to redirect spark output to a file: <https://learn.upgrad.com/v/course/119/question/123151>
5. This link mentioned in resources section of capstone project, was used to get rid of issue faced and added export SPARK\_KAFKA\_VERSION=0.10 : <https://learn.upgrad.com/v/course/119/question/123152>
6. This link mentioned in resources section of capstone project, was used for Spark-HBase configuration and operations: <https://learn.upgrad.com/v/course/119/question/123146>
7. Spark and hive commands are executed via ec2-user.
8. DistanceUtility.java provided was used as is as such with addition of following statements:

At the top of the file: **package com.upgrad.creditcardfrauddetection;**

In DistanceUtilit() method, at the end: **br.close();**

9. Copy FraudAnalysis-0.0.1-SNAPSHOT-jar-with-dependencies.jar and zipCodePosId.csv in /home/ec2-user directory of ec2 instance and run below commands to execute the program:

```
export SPARK_KAFKA_VERSION=0.10
```

```
spark2-submit \  
--class com.upgrad.creditcardfrauddetection.KafkaConsumer \  
--master yarn --deploy-mode client \  
--name CreditCardFraudDetection \  
--conf "spark.app.id=CreditCardFraudDetection spark.driver.memory=12g spark.executor.memory=12g  
spark.executor.instances=2" \  
/home/ec2-user/FraudAnalysis-0.0.1-SNAPSHOT-jar-with-dependencies.jar \  
ec2-54-237-193-160.compute-1.amazonaws.com \  
&> /tmp/Output.txt
```

**NOTE:** Highlighted portion is public DNS for my ec2 instance. Passing it as argument rather than hard coding in the program because it keeps on changing with restart of ec2 instance every time.

10. Run below command to extract card\_transactions\_hbase table into a file:

```
hive -S -e 'use capstone_project; set hive.cli.print.header=true; select * from card_transactions_hbase' | sed  
's/[\\t]/,/g' | sed 's/card_transactions_hbase\//,/g' | sed '/^WARN:/d' > /tmp/card_transactions.csv
```

11. Total records in submitted card\_transactions.csv are: 59367 (excluding header)

**GENUINE: 59150**  
**FRAUD: 217**

12. Original card\_transactions.csv had 53292 records:

**GENUINE: 53210**  
**FRAUD: 82**

```
ec2-user@ip-172-31-91-95:~/capstone_project  
[ec2-user@ip-172-31-91-95 capstone_project]$ ls -lrt card_transactions.csv  
-rw-rw-r-- 1 ec2-user ec2-user 4829520 May 25 05:04 card_transactions.csv  
[ec2-user@ip-172-31-91-95 capstone_project]$ grep GENUINE card_transactions.csv |wc -l  
53210  
[ec2-user@ip-172-31-91-95 capstone_project]$ grep FRAUD card_transactions.csv |wc -l  
82  
[ec2-user@ip-172-31-91-95 capstone_project]$
```

**Task 5:** Create a streaming data processing framework which ingests real-time POS transaction data from Kafka. The transaction data is then validated based on the three rules' parameters (stored in the NoSQL database) discussed above.

**Task 6:** Update the transactions data along with the status (Fraud/Genuine) in the card\_transactions table.

**Task 7:** Store the 'postcode' and 'transaction\_dt' of the current transaction in the look-up table in the NoSQL database.

### ===== Logic Explanation for Task 5 - 6 -7 – Start =====

**FraudAnalysis** is the maven project containing 3 java files namely:

1. **DistanceUtility.java**
2. **KafkaConsumer.java**
3. **CreditCardFraudDetection.java**

#### 1. DistanceUtility.java

It was already provided and used as is. Following statements were added in it in order to compile:

At the top of the file:

```
package com.upgrad.creditcardfrauddetection;
```

In DistanceUtility class, at the end:

```
br.close();
```

This is a Utility class that reads file zipCodePosId.csv and using the same if two post codes are provided as input, it returns distance between two post codes.

The getSpeed method in CreditCardFraudDetection class uses this utility to calculate the distance between 2 post codes. One post code in current transaction record and the other one retrieved from the lookup table which had the latest postcode for a card from last GENUINE transaction. This is one of validation rule for validating a transaction as GENUINE / FRAUD.

Below are the screenshots for DistanceUtility class:

```
package com.upgrad.creditcardfrauddetection;

import java.io.BufferedReader;
import java.io.FileReader;
import java.io.IOException;
import java.util.HashMap;

/**
 * Utility class that reads file zipCodePosId.csv and using same if two zip
 * codes are provided, it returns distances.
 */
class ZipCodeData {
    double lat;
    double lon;
    String city;
    String state_name;
    String postId;

    public ZipCodeData(double lat, double lon, String city, String state_name, String postId) {
        this.lat = lat;
        this.lon = lon;
        this.city = city;
        this.state_name = state_name;
        this.postId = postId;
    }
}

class DistanceUtility {

    HashMap<String, ZipCodeData> zipCodesMap = new HashMap<String, ZipCodeData>();

    /**
     * Initialize zip codes using given file
     *
     * @throws IOException
     * @throws NumberFormatException
     */
    public DistanceUtility() throws NumberFormatException, IOException {

        BufferedReader br = new BufferedReader(new FileReader("zipCodePosId.csv"));

        String line = null;

        while ((line = br.readLine()) != null) {
            String str[] = line.split(",");
            String zipCode = str[0];

            double lat = Double.parseDouble(str[1]);
            double lon = Double.parseDouble(str[2]);
            ;
            String city = str[3];
            String state_name = str[4];
            String postId = str[5];

            ZipCodeData zipCodeData = new ZipCodeData(lat, lon, city, state_name, postId);
        }
    }
}
```

```

        zipCodesMap.put(zipCode, zipCodeData);
    }
    br.close();
}

/**
 *
 * @param zipcode1 - zip code of previous transaction
 * @param zipcode2 - zip code of current transaction
 * @return distance between two zip codes
 */
public double getDistanceViaZipCode(String zipcode1, String zipcode2) {
    ZipCodeData z1 = zipCodesMap.get(zipcode1);
    ZipCodeData z2 = zipCodesMap.get(zipcode2);
    return distance(z1.lat, z1.lon, z2.lat, z2.lon);
}

private double distance(double lat1, double lon1, double lat2, double lon2) {
    double theta = lon1 - lon2;
    double dist = Math.sin(deg2rad(lat1)) * Math.sin(deg2rad(lat2))
        + Math.cos(deg2rad(lat1)) * Math.cos(deg2rad(lat2)) * Math.cos(deg2rad(theta));
    dist = Math.acos(dist);
    dist = rad2deg(dist);
    dist = dist * 60 * 1.1515;
    dist = dist * 1.609344;

    return dist;
}

private double rad2deg(double rad) {
    return rad * 180.0 / Math.PI;
}

private double deg2rad(double deg) {
    return deg * Math.PI / 180.0;
}

}

/*
* //To get the distance between two zipcodes from the main class public class
* PostcodeCalculator { public static void main(String args[]) throws
* NumberFormatException, IOException { DistanceUtility disUtil=new
* DistanceUtility();
*
* System.out.println(disUtil.getDistanceViaZipCode("10001", "10524")); } }
*/

```

## 2. KafkaConsumer.java

This is the main program which connects with Kafka and fetch POS transactions which are in JSON format. Main logic was provided already in resources section of capstone project and used as is with few modifications to invoke logic for Credit Card Fraud Detection.

KafkaConsumer has main method which first setup SparkConf and then invokes JavaStreamingContext. This then setup context with SparkConf and batch duration of 1 second for the incoming D-Streams.

The application subscribes to Kafka stream as consumer for topic - **transactions-topic-verified**. Stream being generated by Kafka is then held by a Collection<String> of topics. It is then converted into JavaDStreams which are further processed to solve given problem statement for Credit Card Fraud Detection.

For the DStreams having Kafka topics, we are interested only in Value part so first fetch Value part by using Map transformation.

Once transactions are obtained, print number of transactions fetched and print all transactions fetched from Kafka topic.

Data from input stream is then mapped to the constructor of CreditCardFraudDetection class by using Map transformation.

After invoking CreditCardFraudDetection class constructor, each record in the JavaDStream becomes an object of type CreditCardFraudDetection and then invokes method in CreditCardFraudDetection class namely FraudDetection in order to classify transaction into GENUINE/FRAUD and update NoSQL DB.

Print Start time, start streaming and wait for termination. Upon receiving Ctrl+C response, gracefully close streaming and print End time.

**Keep on monitoring the output file, few minutes after all transactions are processed probably in 3-4 minutes, Ctrl+C can be pressed to stop the program.**

Below are the screenshots of KafkaConsumer class:

```
package com.upgrad.creditcardfraudetection;

import java.text.SimpleDateFormat;
import java.util.Arrays;
import java.util.Collection;
import java.util.Date;
import java.util.HashMap;
import java.util.Map;

import org.apache.kafka.clients.consumer.ConsumerRecord;
import org.apache.kafka.common.serialization.StringDeserializer;
import org.apache.log4j.Level;
import org.apache.log4j.Logger;
import org.apache.spark.SparkConf;
import org.apache.spark.api.java.JavaRDD;
import org.apache.spark.api.java.function.VoidFunction;
import org.apache.spark.streaming.Durations;
import org.apache.spark.streaming.api.java.JavaDStream;
import org.apache.spark.streaming.api.java.JavaInputDStream;
import org.apache.spark.streaming.api.java.JavaStreamingContext;
import org.apache.spark.streaming.kafka010.ConsumerStrategies;
import org.apache.spark.streaming.kafka010.KafkaUtils;
import org.apache.spark.streaming.kafka010.LocationStrategies;

public class KafkaConsumer {

    /*
     * Setup GROUP_ID in such a way that unique every time
     */
    public static String GROUP_ID = "amitgoelkafkaspark"
        + new SimpleDateFormat("yyyy-MM-dd HH:mm:ss").format(new Date());

    public static void main(String[] args) throws Exception {

        /*
         * Check if 1 argument is passed to the program
         */
        if (args.length != 1) {
            System.out.println("Please enter 1st argument as host server IP");
            return;
        }
        /*
         * Print GROUP_ID for current Kafka stream
         */
        System.out.println("Using group ID : " + GROUP_ID + " for current Kafka stream.");

        /*
         * Set Logger to OFF
         */
        Logger.getLogger("org").setLevel(Level.OFF);
        Logger.getLogger("akka").setLevel(Level.OFF);

        /*
         * Set Spark Configuration object
         */
        SparkConf sparkConf = new SparkConf().setAppName("CreditCardFraudDetection").setMaster("local[*]");
    }
}
```

```

/*
 * Set Streaming Context
 */
JavaStreamingContext jssc = new JavaStreamingContext(sparkConf, Durations.seconds(1));

/*
 * Initialize static variable hostServerIP in CreditCardFraudDetection class.
 * The HBase master and ZooKeeper are dependent on this IP, which keeps on
 * changing with restart of EC2 instance every time.
 */
CreditCardFraudDetection.hostServerIP = args[0];

/*
 * Set parameters for Kafka stream
 */
Map<String, Object> kafkaParams = new HashMap<>();
kafkaParams.put("bootstrap.servers", "100.24.223.181:9092");
kafkaParams.put("key.deserializer", StringDeserializer.class);
kafkaParams.put("value.deserializer", StringDeserializer.class);
kafkaParams.put("group.id", GROUP_ID);
kafkaParams.put("auto.offset.reset", "earliest");
kafkaParams.put("enable.auto.commit", true);

/*
 * Subscribe to Kafka topic
 */
Collection<String> topics = Arrays.asList("transactions-topic-verified");

/*
 * Consume input Kafka Stream and convert into Spark DStreams
 */
JavaInputDStream<ConsumerRecord<String, String>> stream = KafkaUtils.createDirectStream(jssc,
    LocationStrategies.PreferConsistent(),
    ConsumerStrategies.<String, String>Subscribe(topics, kafkaParams));

/*
 * Use map transformation on input stream and get the value part
 */
JavaDStream<String> jds = stream.map(x -> x.value());

/*
 * Print count of records present
 */
jds.foreachRDD(x -> System.out.println("Total Number of Transactions to be Processed : " + x.count() + "\n"));

/*
 * Print all the data so can be used for verification later if needed
 */
jds.foreachRDD(new VoidFunction<JavaRDD<String>>() {
    private static final long serialVersionUID = 1L;

    @Override
    public void call(JavaRDD<String> rdd) {
        rdd.foreach(a -> System.out.println(a));
    }
});

```

```

/*
 * Map data from input stream to CreditCardFraudDetection class constructor
 */
JavaDStream<CreditCardFraudDetection> jds_mapped = jds.map(x -> new CreditCardFraudDetection(x));

/*
 * Call FraudDetection method in CreditCardFraudDetection class
 */
jds_mapped.foreachRDD(new VoidFunction<JavaRDD<CreditCardFraudDetection>>() {
    private static final long serialVersionUID = 1L;

    @Override
    public void call(JavaRDD<CreditCardFraudDetection> rdd) {
        rdd.foreach(x -> x.FraudDetection(x));
    }
});

/*
 * Print current time stamp before starting
 */
System.out.println("\nStart Time : " + new SimpleDateFormat("yyyy-MM-dd HH:mm:ss").format(new Date()) + "\n");

/*
 * Start Spark Streaming
 */
jssc.start();

/*
 * Await Termination to respond to Ctrl+C and gracefully close Spark Streaming
 */
jssc.awaitTermination();

/*
 * Print current time stamp before closing
 */
System.out.println("\nEnd Time : " + new SimpleDateFormat("yyyy-MM-dd HH:mm:ss").format(new Date()) + "\n");

/*
 * Close Spark Streaming
 */
jssc.close();
}
}

```

### 3. CreditCardFraudDetection.java

This class is having business logic to validate incoming transaction against given validation rules and classify into GENUINE/FRAUD. This also updates NoSQL DB with details of current transaction.

This class has a **constructor** which is invoked by main method of KafkaConsumer class and each incoming transaction gets mapped to object of this class. It initializes member variables to values derived after parsing JSON from incoming transaction.

```

public class CreditCardFraudDetection implements java.io.Serializable {

    /*
     * Declare member variables for this class
     */
    private String card_id = null;
    private String member_id = null;
    private double amount = 0.0;
    private String postcode = null;
    private String pos_id = null;
    private Date transaction_dt = null;
    private String status = null;

    private static Admin hBaseAdmin = null;
    public static String hostServerIP = null;

    /*
     * Default constructor defined for this class
     */
    public CreditCardFraudDetection() {
    }

    /*
     * Parameterized constructor for this class. It takes the incoming record from
     * DStream as JSON and parse it. This constructor initializes member variables
     * to the values derived from parsing the JSON.
     */
    public CreditCardFraudDetection(String jsonObj) {

        JSONParser parser = new JSONParser();
        try {

            JSONObject obj = (JSONObject) parser.parse(jsonObj);

            this.card_id = String.valueOf(obj.get("card_id"));
            this.member_id = String.valueOf(obj.get("member_id"));
            this.amount = Double.parseDouble(String.valueOf(obj.get("amount")));
            this.postcode = String.valueOf(obj.get("postcode"));
            this.pos_id = String.valueOf(obj.get("pos_id"));
            this.transaction_dt = new SimpleDateFormat("dd-MM-yyyy HH:mm:ss")
                .parse(String.valueOf(obj.get("transaction_dt")));

        } catch (ParseException e) {
            e.printStackTrace();
        } catch (Exception e) {
            e.printStackTrace();
        }
    }
}

```

This class has a method namely **getHbaseAdmin** for setting up HBase connection. Server IP is passed as command line argument and initialized by main method of KafkaConsumer class. **Hard coding of server IP address is avoided here because public DNS keeps on changing for ec2 instance with every restart.**

```

/*
 * getHbaseAdmin method for setting up the HBase connection. hostServerIP is
 * passed as command line argument and initialized by main method of
 * KafkaConsumer class.
 */
public static Admin getHbaseAdmin() throws IOException {

    Configuration conf = HBaseConfiguration.create();
    conf.setInt("timeout", 120000);
    conf.set("hbase.master", hostServerIP + ":60000");
    conf.set("hbase.zookeeper.quorum", hostServerIP);
    conf.set("hbase.zookeeper.property.clientPort", "2181");
    conf.set("zookeeper.znode.parent", "/hbase");
    try {
        Connection con = ConnectionFactory.createConnection(conf);
        if (hBaseAdmin == null) {
            hBaseAdmin = con.getAdmin();
        }
    } catch (IOException e) {
        e.printStackTrace();
    }

    return hBaseAdmin;
}

```

Main method of CreditCardFraudDetection class is **FraudDetection**. This is the method which is called for every incoming transaction. This method is called by main method of KafkaConsumer class.

It first classifies incoming transaction through classifyTransaction method into GENUINE/FRAUD. It then updates HBase tables by calling updateNoSQLDB method.

It also displays start time when processing of incoming transaction has started and end time when processing of transaction has ended.

```

/*
 * This is the method called by DStream records. It first classify transaction
 * through classifyTransaction method. It then calls updateNoSQLDB method to
 * update data in HBase tables.
 */
public void FraudDetection(CreditCardFraudDetection obj) {
    try {
        System.out.println("\n\n=====");
        /*
         * Print time when transaction processing started
         */
        System.out.println("\nNew Transaction Processing Start Time : "
            + new SimpleDateFormat("yyyy-MM-dd HH:mm:ss").format(new Date()));

        System.out.println("\n===== Classify Current Transaction into GENUINE/FRAUD =====");

        /*
         * classifyTransaction method to validate transaction against various rules and
         * classify current transaction into GENUINE/FRAUD
         */
        classifyTransaction();

        System.out.println("\n===== Update NoSQL Database for Current Transaction =====");

        /*
         * updateNoSQLDB method to update NoSQL DB with card transaction details and if
         * transaction is GENUINE, update lookup table as well
         */
        updateNoSQLDB(obj);

        System.out.println("\n=====");
        /*
         * Print time when transaction processing finished
         */
        System.out.println("\nCurrent Transaction Processing End Time : "
            + new SimpleDateFormat("yyyy-MM-dd HH:mm:ss").format(new Date()));

        System.out.println("\n=====");
    } catch (IOException e) {
        e.printStackTrace();
    }
}

```

Method **classifyTransaction** called by FraudDetection method, apply provided 3 validation rules on each incoming transaction:

1. If credit score of the member is less than 200, then member could be a defaulter and it could be a possible case of fraud.
2. If transaction amount for incoming record/transaction, exceeds UCL (upper control limit) based on last 10 genuine transactions then it could be a possible case of fraud.
3. If speed of travel based on distance between postcode of current and last transaction, is more than 1km in 4sec = 0.25 km/sec then it could be a possible case of fraud

This method first finds out UCL for Card ID from HBase lookup table by calling method getUCL and displays UCL and current transaction amount. Then method finds out credit score of member owning Card ID from HBase lookup table by calling method getScore and displays credit score. Then method calls getSpeed method to find out speed at which member could have travelled based on distance between post code for current transaction and post code for last GENUINE transaction.

Explanation so far is shown in screenshot below:

```
/*
 * Classify incoming transaction at POS as GENUINE or FRAUD
 */
private void classifyTransaction() throws IOException {

    /*
     * Lets apply provided rules on each transaction:
     *
     * 1. If member score is less than 200, then member could be a defaulter and
     * possible case of fraud.
     *
     * 2. If transaction amount for current transaction, exceeds UCL (upper control
     * limit) based on last 10 genuine transactions then possible case of fraud.
     *
     * 3. If speed of travel based on distance between postcode of current and last
     * transaction, is more than 1km in 4sec = 0.25 km/sec then possible case of
     * fraud.
     */

    double card_ucl = 0;
    int card_score = 0;
    double speed = 0;

    /*
     * Print Card ID of current transaction
     */
    System.out.println("\nCard ID of current transaction is : " + this.getCard_id());

    /*
     * Get UCL for Card ID
     */
    card_ucl = getUCL(this.getCard_id());
    System.out.println("\nUCL (Upper Control Limit) for Card ID of current transaction is : " + card_ucl);

    /*
     * Print current transaction amount
     */
    System.out.println("\nCurrent Transaction Amount is : " + this.getAmount());

    /*
     * Get score for Card ID
     */
    card_score = getScore(this.getCard_id());
    System.out.println("\nCredit score for Card ID of current transaction is : " + card_score);

    /*
     * Get speed of current transaction in relation to last transaction
     */
    speed = getSpeed();
```

Method then checks if credit score  $\geq 200$  and amount of current transaction is  $\leq$  UCL and speed is  $\leq 0.25$  and if so then status of the transaction is classified as ‘GENUINE’ as it has passed all 3 validation rules. Details are displayed accordingly.

If any of the validation rules fail then status of the transaction is classified as ‘FRAUD’ and failed validation rule is displayed for the information.

Explanation so far is shown in screenshot below:

```

/*
 * All 3 rules need to be passed for card transaction status to be GENUINE
 */
if ((card_score >= 200) && (this.getAmount() <= card_ucl) && (speed <= 0.25)) {
    this.status = "GENUINE";
    System.out.println("\nCurrent Card Transaction has passed Credit Score Validation rule : Score : "
        + card_score + " >= 200");
    System.out.println(
        "\nCurrent Card Transaction has passed UCL (Upper Control Limit) Validation rule : Amount : "
        + this.getAmount() + " <= " + card_ucl);
    System.out.println(
        "\nCurrent Card Transaction has passed Zip Code Distance Validation rule : Speed (km/sec) : "
        + speed + " <= 0.25");
    System.out.println("\nCurrent Card Transaction has passed all 3 Validation rules and so status is : "
        + this.getStatus());
} else {
    this.status = "FRAUD";
    if (card_score < 200) {
        System.out.println("\nCurrent Card Transaction did not pass through Credit Score Validation rule");
    }
    if (speed > 0.25) {
        System.out.println("\nCurrent Card Transaction did not pass through Zip Code Distance Validation rule");
    }
    if (this.getAmount() > card_ucl) {
        System.out.println(
            "\nCurrent Card Transaction did not pass through UCL (Upper Control Limit) Validation rule");
    }
    System.out.println("\nCurrent Card Transaction status is : " + this.getStatus());
}
}

```

Method **getUCL** called by classifyTransaction method, first establishes HBase connection if not established already by calling getHbaseAdmin method. It then fetches UCL for input Card ID from lookup\_data\_hive lookup HBase table. This method code is shown in screenshot below:

```

/*
 * getUCL method to look up upper control limit for input Card ID from
 * lookup_data_hive HBASE table
 */
public static Double getUCL(String cardID) throws IOException {

    if (cardID == null) {
        System.out.println("\nCard ID is not present in data received from Kafka. Kindly check Kafka stream");
        return -1d;
    } else {

        /*
         * Call getHbaseAdmin() method only if connection is not already established
         */
        if (hBaseAdmin == null) {
            hBaseAdmin = getHbaseAdmin();
        }

        /*
         * Setup table name, get data for input card id and get ucl from result set. ucl
         * belongs to lookup_card_family column family in look up table
         * lookup_data_hive.
         */
        String table = "lookup_data_hive";
        double val = 0;
        try {
            Get cardId = new Get(Bytes.toBytes(cardID));
            Table htable = hBaseAdmin.getConnection().getTable(TableName.valueOf(table));
            Result result = htable.get(cardId);
            byte[] value = result.getValue(Bytes.toBytes("lookup_card_family"), Bytes.toBytes("ucl"));
            if (value != null) {
                val = Double.parseDouble(Bytes.toString(value));
            } else {
                val = 0d;
            }
        } catch (Exception e) {
            e.printStackTrace();
        }
        return val;
    }
}

```

Method **getScore** called by classifyTransaction method, first establishes HBase connection if not established already by calling getHbaseAdmin method. It then fetches credit score for input Card ID from lookup\_data\_hive lookup HBase table. This method code is shown in screenshot below:

```

/*
 * getScore method to look up score for input Card ID from lookup_data_hive
 * HBASE table
 */
public static Integer getScore(String cardID) throws IOException {
    if (cardID == null) {
        System.out.println("\nCard ID is not present in data received from Kafka. Kindly check Kafka stream");
        return -1;
    } else {
        /*
         * Call getHbaseAdmin() method only if connection is not already established
         */
        if (hBaseAdmin == null) {
            hBaseAdmin = getHbaseAdmin();
        }

        /*
         * Setup table name, get data for input card id and get score from result set.
         * score belongs to lookup_card_family column family in look up table
         * lookup_data_hive.
         */
        String table = "lookup_data_hive";
        int val = 0;
        try {
            Get cardId = new Get(Bytes.toBytes(cardID));
            Table htable = hBaseAdmin.getConnection().getTable(TableName.valueOf(table));
            Result result = htable.get(cardId);
            byte[] value = result.getValue(Bytes.toBytes("lookup_card_family"), Bytes.toBytes("score"));
            if (value != null) {
                val = Integer.parseInt(Bytes.toString(value));
            } else
                val = 0;

        } catch (Exception e) {
            e.printStackTrace();
        }
        return val;
    }
}

```

Method **getSpeed** called by `classifyTransaction` method, first creates an object of `DistanceUtility` class. It then fetches post code of last genuine transaction for input Card ID from HBase lookup table by calling `getPostCode` method and display post code of both last genuine and current transaction. It then fetches transaction date of last genuine transaction for input Card ID from HBase lookup table by calling `getTransactionDate` method and display transaction date of both last genuine and current transaction. It then calls `getDistanceViaZipCode` method on object of `DistanceUtility` class by passing post code of last genuine and current transaction. This call returns distance between 2 post codes in kilo meters (Km). Method then calculates time difference between transaction date of last genuine and current transaction which comes out in milli second and divide by 1000 to get in seconds. Speed is then calculated by dividing distance with time difference.  $\text{Speed} = \frac{\text{Distance}}{\text{Time}}$ . Distance between post code of last genuine and current transaction is divided by time difference between transaction date of last genuine and current transaction to obtain speed which comes out in km/sec. This is then returned to calling method.

Explanation so far is shown in screenshots below:

```

/*
 * getSpeed method to calculate speed in km/sec based on distance between post
 * code of current and last transaction. This method makes of DistanceUtility
 * class provided.
 */
private double getSpeed() {

    double distance = 0;
    long timeDifference = 0L;
    double speed = 0;

    try {

        /*
         * Create an object of DistanceUtility class
         */
        DistanceUtility distUtil = new DistanceUtility();

        /*
         * Get last post code of Card ID from lookup table in HBase
         */
        String lastPostCode = getPostCode(this.getCard_id());
        System.out.println("\nLast Post Code for Card ID of current transaction is : " + lastPostCode);

        /*
         * Print post code of current transaction
         */
        System.out.println("\nCurrent Post Code for Card ID of current transaction is : " + this.getPostcode());

        /*
         * Get last transaction_dt of Card ID from Lookup table in HBase
         */
        Date lastTransactionDt = getTransactionDate(this.getCard_id());
        System.out.println("\nLast Transaction Date for Card ID of current transaction is : "
            + new SimpleDateFormat("yyyy-MM-dd HH:mm:ss").format(lastTransactionDt));

        /*
         * Print transaction date of current transaction
         */
        System.out.println("\nCurrent Transaction Date for Card ID of current transaction is : "
            + new SimpleDateFormat("yyyy-MM-dd HH:mm:ss").format(this.getTransaction_dt()));

        /*
         * Get distance between 2 post code, calculated by following method in
         * DistanceUtility class. Distance is in kilometers.
         */
        distance = distUtil.getDistanceViaZipCode(lastPostCode, this.getPostcode());
        System.out
            .println("\nDistance (returned by DistanceUtility) between current and last postcode (in km) is : "
                + distance);
    }
}

```

```

/*
 * Time difference between 2 transaction dates is calculated as below. The
 * difference is returned in milliseconds so dividing by 1000 to get time in
 * seconds. For many transactions, the incoming transaction date is less than
 * last transaction date stored in lookup table so using abs method.
 */
timeDifference = (java.lang.Math.abs(this.getTransaction_dt().getTime() - lastTransactionDt.getTime()))
    / 1000;
System.out.println("\nAbsolute Time Difference between current and last transaction (in seconds) is : "
    + timeDifference);

/*
 * Speed is being calculated in km/sec
 *
 * At one instance, time difference was zero (2018-06-21 15:29:59) and last and
 * current post code (55311) were same, distance utility returned non zero value
 * (9.493073054631141E-5) which is quite small as such but still a non zero
 * value was returned, due to which speed came out to be infinity for card_id =
 * '4838289241690162'. So added this check if last and current post code are
 * same then consider speed as 0.
 */

if (lastPostCode.equals(this.getPostcode())) {
    System.out.println(
        "\nSince Last Post Code and Current Post Code are same, so distance is 0 and so considering speed as 0");
    speed = 0;
} else {
    speed = distance / timeDifference;
}

System.out.println("\nSpeed (Distance / Time) : " + speed + " km/second ");

} catch (Exception e) {
    e.printStackTrace();
}

return speed;
}

```

=====

If post code of last genuine transaction and current transaction is same then distance is considered as 0. There is a problem with DistanceUtility provided that it returns non-zero value at times when post code of last genuine transaction and current transaction is same. That is not ideal so considering distance as 0 in such cases to avoid this problem.

I have provided Output.txt as one of the artifact. Put this in a directory on ec2 instance and run this unix command:

**grep "Distance (returned by DistanceUtility) between current and last postcode (in km) is : .\*E-.\*" Output.txt**

Below output will be returned by the command:

Distance (returned by DistanceUtility) between current and last postcode (in km) is : 9.493073054631141E-5  
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 9.493073054631141E-5  
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 9.493073054631141E-5  
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 1.3425232662457948E-4  
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 1.3425232662457948E-4  
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 9.493073054631141E-5

These are all cases where DistanceUtility returned non zero value while should have returned zero because post code of last genuine and current transaction is same.

Below are some more details fetched from Output.txt for these cases:

Card ID of current transaction is : 4838289241690162

Last Post Code for Card ID of current transaction is : 55311

Current Post Code for Card ID of current transaction is : 55311

Distance (returned by DistanceUtility) between current and last postcode (in km) is : 9.493073054631141E-5

Card ID of current transaction is : 371829395919892

Last Post Code for Card ID of current transaction is : 60929

Current Post Code for Card ID of current transaction is : 60929

Distance (returned by DistanceUtility) between current and last postcode (in km) is : 9.493073054631141E-5

Card ID of current transaction is : 4750699680073601

Last Post Code for Card ID of current transaction is : 88339

Current Post Code for Card ID of current transaction is : 88339

Distance (returned by DistanceUtility) between current and last postcode (in km) is : 9.493073054631141E-5

Card ID of current transaction is : 4552141762363568

Last Post Code for Card ID of current transaction is : 55735

Current Post Code for Card ID of current transaction is : 55735

Distance (returned by DistanceUtility) between current and last postcode (in km) is : 1.3425232662457948E-4

Card ID of current transaction is : 344997383137673

Last Post Code for Card ID of current transaction is : 89043

Current Post Code for Card ID of current transaction is : 89043

Distance (returned by DistanceUtility) between current and last postcode (in km) is : 1.3425232662457948E-4

Card ID of current transaction is : 344250415187512

Last Post Code for Card ID of current transaction is : 97015

Current Post Code for Card ID of current transaction is : 97015

Distance (returned by DistanceUtility) between current and last postcode (in km) is : 9.493073054631141E-5

=====

Method **getPostCode** called by getSpeed method, first establishes HBase connection if not established already by calling getHbaseAdmin method. It then fetches post code for input Card ID from lookup\_data\_hive lookup HBase table. This method code is shown in screenshot below:

```
/*
 * getPostCode method to look up last post code for input Card ID from
 * lookup_data_hive HBASE table
 */
public static String getPostCode(String cardID) throws IOException {

    if (cardID == null) {
        System.out.println("\nCard ID is not present in data received from Kafka. Kindly check Kafka stream");
        return null;
    } else {

        /*
         * Call getHbaseAdmin() method only if connection is not already established
         */
        if (hBaseAdmin == null) {
            hBaseAdmin = getHbaseAdmin();
        }

        /*
         * Setup table name, get data for input card id and get postcode from result
         * set. postcode belongs to lookup_transaction_family column family in look up
         * table lookup_data_hive.
         */
        String table = "lookup_data_hive";
        String val = null;
        try {
            Get cardId = new Get(Bytes.toBytes(cardID));
            Table htable = hBaseAdmin.getConnection().getTable(TableName.valueOf(table));
            Result result = htable.get(cardId);
            byte[] value = result.getValue(Bytes.toBytes("lookup_transaction_family"), Bytes.toBytes("postcode"));
            if (value != null) {
                val = Bytes.toString(value);
            } else
                val = null;
        } catch (Exception e) {
            e.printStackTrace();
        }
        return val;
    }
}
```

Method **getTransactionDate** called by getSpeed method, first establishes HBase connection if not established already by calling getHbaseAdmin method. It then fetches transaction date for input Card ID from lookup\_data\_hive lookup HBase table. This method code is shown in screenshot below:

```
/*
 * getTransactionDate method to look up last transaction date for input Card ID
 * from lookup_data_hive HBASE table
 */
public static Date getTransactionDate(String cardID) throws IOException {

    if (cardID == null) {
        System.out.println("\nCard ID is not present in data received from Kafka. Kindly check Kafka stream");
        return null;
    } else {

        /*
         * Call getHbaseAdmin() method only if connection is not already established
         */
        if (hBaseAdmin == null) {
            hBaseAdmin = getHbaseAdmin();
        }

        /*
         * Setup table name, get data for input card id and get transaction_dt from
         * result set. transaction_dt belongs to lookup_transaction_family column family
         * in look up table lookup_data_hive.
         */
        String table = "lookup_data_hive";
        Date val = null;
        try {
            Get cardId = new Get(Bytes.toBytes(cardID));
            Table htable = hBaseAdmin.getConnection().getTable(TableName.valueOf(table));
            Result result = htable.get(cardId);
            byte[] value = result.getValue(Bytes.toBytes("lookup_transaction_family"),
                Bytes.toBytes("transaction_dt"));
            if (value != null) {
                val = new SimpleDateFormat("yyyy-MM-dd HH:mm:ss").parse(Bytes.toString(value));
            } else
                val = null;
        } catch (Exception e) {
            e.printStackTrace();
        }
        return val;
    }
}
```

**Method updateNoSQLDB called by FraudDetection method in CreditCardFraudDetection class accomplishes Task 6 and Task 7 from the given problem statement.**

It performs 2 activities:

1. It inserts incoming card transaction details in card\_transactions\_hive HBase table. **(Task 6)**

It first establishes HBase connection if not established already by calling getHbaseAdmin method. It converts transaction date of incoming transaction in a suitable format for HBase table. Row key of HBase table is populated using randomUUID() so method uses randomUUID() as row key for the HBase table. New Put object is created using row key and then other columns are added to it like card\_id, member\_id, amount, postcode, pos\_id, transaction\_dt and status. These are all then saved down in card\_transactions\_hive HBase table. If by chance, HBase table does not exist then appropriate message gets displayed accordingly.

2. It also updates lookup\_data\_hive lookup HBase table with current post code and transaction date if and only if the status of current transaction is GENUINE. **(Task 7)**

It first establishes HBase connection if not established already by calling getHbaseAdmin method. Card ID is used as row key for lookup HBase table. New Put object is created using row key and then transaction date (in a suitable format for HBase table) and post code are added to it. These are all then saved down in lookup\_data\_hive lookup HBase table. If by chance, HBase lookup table does not exist then appropriate message gets displayed accordingly.

Explanation so far is shown in screenshots below:

```
/*
 * updateNoSQLDB method performs 2 activities:
 *
 * 1. It inserts incoming card transaction details in card_transactions_hive
 * HBASE table.
 *
 * 2. It also updates lookup_data_hive HBASE table with current post code and
 * transaction date if and only if the status of current transaction is GENUINE.
 */
public static void updateNoSQLDB(CreditCardFraudDetection transactionData) throws IOException {

    try {

        /*
         * Call getHbaseAdmin() method only if connection is not already established
         */
        if (hBaseAdmin == null) {
            hBaseAdmin = getHbaseAdmin();
        }

        /*
         * Print Card ID of current transaction
         */
        System.out.println("\nCard ID of current transaction is : " + transactionData.getCard_id());

        /*
         * Convert incoming transaction date to specified format
         */
        DateFormat dateFormat = new SimpleDateFormat("yyyy-MM-dd HH:mm:ss");
        String val = dateFormat.format(transactionData.getTransaction_dt());

        Double Amount = transactionData.getAmount();

        String transactionTable = "card_transactions_hive";

        /*
         * Check if table exists
         */
        if (hBaseAdmin.tableExists(TableName.valueOf(transactionTable))) {
            Table htable = hBaseAdmin.getConnection().getTable(TableName.valueOf(transactionTable));

            /*
             * Use randomUUID as row key
             */

            UUID transactionID = UUID.randomUUID();

            Put p = new Put(Bytes.toBytes(transactionID.toString()));

```

```

/*
 * Add column values for each column
 */
p.addColumn(Bytes.toBytes("card_transactions_family"), Bytes.toBytes("card_id"),
           Bytes.toBytes(transactionData.getCard_id()));
p.addColumn(Bytes.toBytes("card_transactions_family"), Bytes.toBytes("member_id"),
           Bytes.toBytes(transactionData.getMember_id()));
p.addColumn(Bytes.toBytes("card_transactions_family"), Bytes.toBytes("amount"),
           Bytes.toBytes(Amount.toString()));
p.addColumn(Bytes.toBytes("card_transactions_family"), Bytes.toBytes("postcode"),
           Bytes.toBytes(transactionData.getPostcode()));
p.addColumn(Bytes.toBytes("card_transactions_family"), Bytes.toBytes("pos_id"),
           Bytes.toBytes(transactionData.getPos_id()));
p.addColumn(Bytes.toBytes("card_transactions_family"), Bytes.toBytes("transaction_dt"),
           Bytes.toBytes(val));
p.addColumn(Bytes.toBytes("card_transactions_family"), Bytes.toBytes("status"),
           Bytes.toBytes(transactionData.getStatus()));

htable.put(p);
System.out.println("\nCurrent Transaction Details are populated in Card Transactions HBase Table : " +
    + transactionTable);
} else {
    System.out.println("\nHBase Table named : " + transactionTable + " : does not exist");
}

} catch (Exception e) {
    e.printStackTrace();
}

/*
 * Check if status is Genuine to update lookup_data_hive HBase table
 */
if (transactionData.getStatus().equals("GENUINE")) {

    System.out.println("\nCurrent Card Transaction status is GENUINE so updating lookup table");

    try {

        /*
         * Call getHbaseAdmin() method only if connection is not already established
         */
        if (hBaseAdmin == null) {
            hBaseAdmin = getHbaseAdmin();
        }

        /*
         * Convert incoming transaction date to specified format
         */
        DateFormat dateFormat = new SimpleDateFormat("yyyy-MM-dd HH:mm:ss");
        String val = dateFormat.format(transactionData.getTransaction_dt());

        String lookupTable = "lookup_data_hive";

        /*
         * Check if table exists
         */
        if (hBaseAdmin.tableExists(TableName.valueOf(lookupTable))) {
            Table htable = hBaseAdmin.getConnection().getTable(TableName.valueOf(lookupTable));

```

```

/*
 * Instantiate Put class to Card ID row key in lookup_data_hive table
 */
Put p = new Put(Bytes.toBytes(transactionData.getCard_id()));

/*
 * Add column values for postcode and transaction_dt
 */
p.addColumn(Bytes.toBytes("lookup_transaction_family"), Bytes.toBytes("postcode"),
    Bytes.toBytes(transactionData.getPostcode()));
p.addColumn(Bytes.toBytes("lookup_transaction_family"), Bytes.toBytes("transaction_dt"),
    Bytes.toBytes(val));

/*
 * Save put instance to HTable.
 */
htable.put(p);
System.out.println(
    "\nPostcode and Transaction Date updated for Card ID of current transaction in lookup HBase table : "
    + lookupTable);

} else {
    System.out.println("\nHBase Lookup Table named : " + lookupTable + " : does not exist");
}
} catch (Exception e) {
    e.printStackTrace();
}
} else {
    System.out.println("\nCurrent Card Transaction status is FRAUD so not updating lookup table");
}

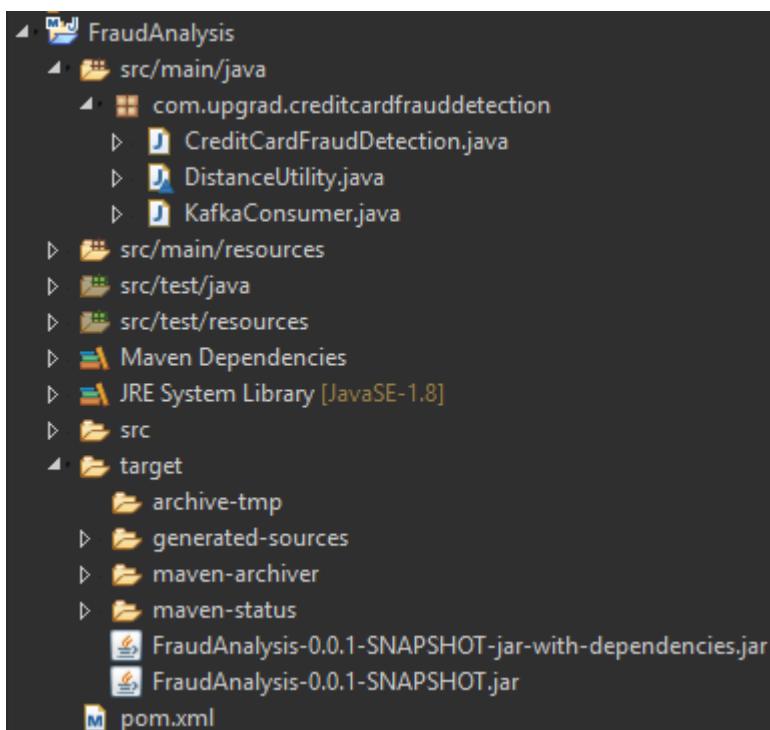
}

```

## ===== Logic Explanation for Task 5 - 6 -7 – END =====

## ===== Code Execution – Start =====

FraudAnalysis.zip has [FraudAnalysis-0.0.1-SNAPSHOT-jar-with-dependencies.jar](#) as well as maven project. Either provided jar can be used or maven project can be imported in eclipse and fresh jar can be generated in eclipse. After project is imported, ensure java build path is pointing to jdk 1.8 and project compliance is also set to 1.8. Do maven clean followed by maven install. It will generate fresh jar with same name.



**Copy FraudAnalysis-0.0.1-SNAPSHOT-jar-with-dependencies.jar and zipCodePosId.csv in /home/ec2-user directory of ec2 instance and run below commands to execute the program:**

```
export SPARK_KAFKA_VERSION=0.10
```

```
spark2-submit \
--class com.upgrad.creditcardfrauddetection.KafkaConsumer \
--master yarn --deploy-mode client \
--name CreditCardFraudDetection \
--conf "spark.app.id=CreditCardFraudDetection spark.driver.memory=12g spark.executor.memory=12g \
spark.executor.instances=2" \
/home/ec2-user/FraudAnalysis-0.0.1-SNAPSHOT-jar-with-dependencies.jar \
ec2-54-237-193-160.compute-1.amazonaws.com \
&> /tmp/Output.txt
```

```
[ec2-user@ip-172-31-91-95:~]
[ec2-user@ip-172-31-91-95 ~]$ export SPARK_KAFKA_VERSION=0.10
[ec2-user@ip-172-31-91-95 ~]$ spark2-submit \
> --class com.upgrad.creditcardfrauddetection.KafkaConsumer \
> --master yarn --deploy-mode client \
> --name CreditCardFraudDetection \
> --conf "spark.app.id=CreditCardFraudDetection spark.driver.memory=12g spark.executor.memory=12g spark.executor.instances=2" \
> /home/ec2-user/FraudAnalysis-0.0.1-SNAPSHOT-jar-with-dependencies.jar \
> ec2-54-237-193-160.compute-1.amazonaws.com \
> &> /tmp/Output.txt
```

**Open another session for ec2 instance and run:**

```
tail -f /tmp/Output.txt
```

```
[ec2-user@ip-172-31-91-95:tmp]
```

```
Current Card Transaction has passed Credit Score Validation rule : Score : 683 >= 200
Current Card Transaction has passed UCL (Upper Control Limit) Validation rule : Amount : 6880309.0 <= 1.0869782716383055E7
Current Card Transaction has passed Zip Code Distance Validation rule : Speed (km/sec) : 1.3972584550746266E-5 <= 0.25
Current Card Transaction has passed all 3 Validation rules and so status is : GENUINE
===== Update NoSQL Database for Current Transaction =====
Card ID of current transaction is : 5391723993945313
Current Transaction Details are populated in Card Transactions HBase Table : card_transactions_hive
Current Card Transaction status is GENUINE so updating lookup table
Postcode and Transaction Date updated for Card ID of current transaction in lookup HBase table : lookup_data_hive
=====
Current Transaction Processing End Time : 2019-06-01 11:15:53
=====
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
Total Number of Transactions to be Processed : 0
```

**Once all transactions are processed, press Ctrl+C on the session where spark command was executed.**

```
[ec2-user@ip-172-31-91-95:~]
```

```
[ec2-user@ip-172-31-91-95 ~]$ export SPARK_KAFKA_VERSION=0.10
[ec2-user@ip-172-31-91-95 ~]$ spark2-submit \
> --class com.upgrad.creditcardfrauddetection.KafkaConsumer \
> --master yarn --deploy-mode client \
> --name CreditCardFraudDetection \
> --conf "spark.app.id=CreditCardFraudDetection spark.driver.memory=12g spark.executor.memory=12g spark.executor.instances=2" \
> /home/ec2-user/FraudAnalysis-0.0.1-SNAPSHOT-jar-with-dependencies.jar \
> ec2-54-237-193-160.compute-1.amazonaws.com \
> t> /tmp/Output.txt
^C[ec2-user@ip-172-31-91-95 ~]$
```

**Below is the screenshot of other session of ec2 instance after program has ended.**

 ec2-user@ip-172-31-91-95:/tmp

End Time : 2019-06-01 11:17:22

**Run below command to extract card\_transactions\_hbase table into a file:**

```
hive -S -e 'use capstone_project; set hive.cli.print.header=true; select * from card_transactions_hbase' | sed 's/[\t]//g' | sed 's/card_transactions_hbase\//g' | sed '/^WARN:/d' > /tmp/card_transactions.csv
```

```
[ec2-user@ip-172-31-91-95 ~]$ hive -S -e 'use capstone_project; set hive.cli.print.header=true; select * from card_transactions_hbase' | sed 's/[\t]//g' | sed 's/card_transactions_hbase\//g' | sed '/^WARN:/d' > /tmp/card_transactions.csv
Java HotSpot(TM) 64-Bit Server VM warning: ignoring option MaxPermSize=612M; support was removed in 8.0
Java HotSpot(TM) 64-Bit Server VM warning: Using incremental CMS is deprecated and will likely be removed in a future release
Java HotSpot(TM) 64-Bit Server VM warning: ignoring option MaxPermSize=612M; support was removed in 8.0
[ec2-user@ip-172-31-91-95 ~]$
```

**Below are the files created in /tmp directory**

```
[ec2-user@ip-172-31-91-95 /tmp]$ ls -lrt card_transactions.csv Output.txt
-rw-rw-r-- 1 ec2-user ec2-user 13157831 Jun  1 11:17 Output.txt
-rw-rw-r-- 1 ec2-user ec2-user 7693571 Jun  1 11:25 card_transactions.csv
[ec2-user@ip-172-31-91-95 /tmp]$
```

**Below are no of transactions classified into GENUINE and FRAUD in final output card\_transactions.csv**

**Total: 59367**

**GENUINE: 59150**

**FRAUD: 217**

```
[ec2-user@ip-172-31-91-95 /tmp]$ grep GENUINE card_transactions.csv | wc -l
59150
[ec2-user@ip-172-31-91-95 /tmp]$ grep FRAUD card_transactions.csv | wc -l
217
[ec2-user@ip-172-31-91-95 /tmp]$
```

**Below are no of transactions provided in original card\_transactions.csv**

**Total: 53292**

**GENUINE: 53210**

**FRAUD: 82**

```
[ec2-user@ip-172-31-91-95 ~/capstone_project]$ ls -lrt card_transactions.csv
-rw-rw-r-- 1 ec2-user ec2-user 4829520 May 25 05:04 card_transactions.csv
[ec2-user@ip-172-31-91-95 capstone_project]$ grep GENUINE card_transactions.csv |wc -l
53210
[ec2-user@ip-172-31-91-95 capstone_project]$ grep FRAUD card_transactions.csv |wc -l
82
[ec2-user@ip-172-31-91-95 capstone_project]$
```

**===== Code Execution – END =====**

===== Output – Start =====

**Below are the screenshots from HBase after spark program was over:**

**Count of records in lookup\_data\_hive lookup table**

```
[ec2-user@ip-172-31-91-95 ~]$ hbase shell
Java HotSpot(TM) 64-Bit Server VM warning: Using incremental CMS is deprecated and will likely be removed in a future release
19/06/01 11:19:13 INFO Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.15.0, rUnknown, Thu May 24 04:27:47 PDT 2018

hbase(main):001:0> count 'lookup_data_hive'
999 row(s) in 0.5010 seconds

=> 999
hbase(main):002:0> █
```

**Count of records in card\_transactions\_hive table**

[ec2-user@ip-172-31-91-95:~]

```
hbase(main):001:0> count 'card_transactions_hive'
Current count: 1000, row: 048516cc-b769-4cf8-bf42-cdc613d99982
Current count: 2000, row: 08c5579d-d6c8-4011-8928-76ddf315e899
Current count: 3000, row: 0cea0bfe-b105-455c-a5e1-3c9e0d165657
Current count: 4000, row: 10f5ea47-2985-4dd2-ae83-d0ce5eaffec1
Current count: 5000, row: 155dfa6e-01b8-4e80-8114-b6a0d450301e
Current count: 6000, row: 19d5e179-81e4-4a96-9118-44e974a6f0b3
Current count: 7000, row: 1dd8a85ff-0566-441b-bd10-b63c07fd4b98
Current count: 8000, row: 22406713-fcdb-4b6d-bf93-621dc697c4c4
Current count: 9000, row: 2688d737-eda4-4636-b73b-0ecaa517ad0d
Current count: 10000, row: 2acf7f89-9a47-4fbc-b054-c96161dac4f8
Current count: 11000, row: 2f69d859-1bf3-4d6e-b0e1-bc86ab32d0b7
Current count: 12000, row: 338a91d2-852f-4187-86ad-af47af5f13a4
Current count: 13000, row: 37d046d8-7c74-4337-aeb4-a0988e06f3d0
Current count: 14000, row: 3c3922e0-2ecd-4419-b31d-62d16ad998a3
Current count: 15000, row: 40a6b2fb-255c-4540-a5b0-bf9361e1ce6e
Current count: 16000, row: 44f7e2f6-2e6b-4ada-9b57-97d558elafc5
Current count: 17000, row: 492513ce-76e4-4d30-9ac7-32cb449f4f10
Current count: 18000, row: 4da0a031-e2a2-4234-9a04-fd8a0bf132ca
Current count: 19000, row: 51d191ee-fde4-4bad-ac0e-8bf061b682b2
Current count: 20000, row: 56494ff8-5ba9-43cd-a223-a53170c07ae8
Current count: 21000, row: 5a961c14-019d-4d25-b8d3-f927929b5f6d
Current count: 22000, row: 5f06439c-51f9-4da7-b51d-563868da8b5d
Current count: 23000, row: 63493804-2434-4d16-b978-2671b9a29c2b
Current count: 24000, row: 67975634-3bcc-4b74-92af-ea8d0006a25e
Current count: 25000, row: 6bd95cfa-b96b-465d-a792-fd17b2c9ed67
Current count: 26000, row: 7025b63f-c3d9-4d89-b02d-6641b34eedf5
Current count: 27000, row: 7445a166-a422-4203-8695-5a3c8f25090f
Current count: 28000, row: 788de720-7f17-4040-8bf4-6078cd54b78c
Current count: 29000, row: 7cee5c9b-70bd-4099-a632-c2962dbc5206
Current count: 30000, row: 8122c635-db39-4610-acla-cfb865669e92
Current count: 31000, row: 856966d4-c697-4635-a9c2-14c98c7clab3
Current count: 32000, row: 89f38c2a-a752-4f23-9b03-18021418cab
Current count: 33000, row: 8e66f3c7-61d9-4995-bd6f-eb36f150653e
Current count: 34000, row: 92c9e120-f7f5-4511-8338-37e64bclea12
Current count: 35000, row: 971bbd7e-feee-4d23-85e6-ed5a360fe769
Current count: 36000, row: 9b7b29a3-93ea-4206-ac44-1134886a5b0f
Current count: 37000, row: 9fdd08fc-5024-4703-b81a-4e8a7dc232c3
Current count: 38000, row: a40921fb-fc8c-47e6-a140-8fb027e15a2
Current count: 39000, row: a815bb5d-d856-418b-a6a5-9a0af01cf499
Current count: 40000, row: ac73e814-be68-47a0-8e17-a56aa416b4aa
Current count: 41000, row: b0ella94-cd8f-4e88-9765-0edcecc278e0
Current count: 42000, row: b5370f72-caf6-408d-99c0-a2f886fd533c
Current count: 43000, row: b9131fa3-1120-41d8-a08c-4a948d475d48
Current count: 44000, row: bd7f69fb-afdb-4aa9-bb02-1002f41bc9cb
Current count: 45000, row: c213bb6e-00eb-4cf1-9a0c-712e6c49af80
Current count: 46000, row: c652b30b-6af6-4040-9e3d-294d984c8f81
Current count: 47000, row: ca8f7a3b-032a-4ed3-91c2-bdc30fb88606
Current count: 48000, row: cf1101ca-b289-4598-a79e-fa70ce5ccb61
Current count: 49000, row: d3453dbd-499a-434f-a7a0-16166320289b
Current count: 50000, row: d7837c76-5aaa-4211-a6f8-5b1782c29ba2
Current count: 51000, row: dbed1lab-aalc-421f-8748-90385bef2006
Current count: 52000, row: e050f84a-f79a-40e7-b8c3-ad8f213b6fca
Current count: 53000, row: e4c0e75c-8b59-4af6-9966-f4bc19bd0850
Current count: 54000, row: e9008eb4-df31-4aac-8152-8664820fde47
Current count: 55000, row: ed3d084f-caf5-49a9-9455-aed812e99b7c
Current count: 56000, row: f16f04e2-79f6-492c-bbaa-83158489f599
Current count: 57000, row: f5b59b60-081a-4466-8151-alfd87502013
Current count: 58000, row: falcba61-fc9b-4ffd-8da3-ac07e04399d0
Current count: 59000, row: fe8b817a-1cb0-44ba-b2c7-49ab94043afa
59367 row(s) in 3.5890 seconds
```

=> 59367

Some records from `lookup_data_hive` lookup table. I had created 2 column families in it namely `lookup_card_family` and `lookup_transaction_family`. Column family `lookup_card_family` has columns `uc1` and `score` and can store up to 1 VERSION. Column family `lookup_transaction_family` has columns `postcode` and `transaction_dt` and can store up to 10 VERSIONS. Please refer to `LogicMid.pdf` for more details.

**Below are the screenshots from Hive after spark program was over:**

Below screenshot shows some records from card\_transactions\_hbase table

```
[ec2-user@ip-172-31-91-95 ~]$ hive
Java HotSpot(TM) 64-Bit Server VM warning: ignoring option MaxPermSize=512M; support was removed in 8.0
Java HotSpot(TM) 64-Bit Server VM warning: Using incremental CMS is deprecated and will likely be removed in a future release
Java HotSpot(TM) 64-Bit Server VM warning: ignoring option MaxPermSize=512M; support was removed in 8.0

Logging initialized using configuration in jar:file:/opt/cloudera/parcels/CDH-5.15.0-1.cdh5.15.0.p0.21/jars/hive-common-1.1.0-cdh5.15.0.jar!/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> use capstone_project;
OK
Time taken: 1.757 seconds
hive> select * from card_transactions_hbase limit 20;
OK
00000d47-75fb-4b26-9b24-2c7b87db6ce2 5175055180166201 689714705796007 167127.0 45764 486276156459636 2017-11-27 08:56:20 GENUINE
0003e0a0-5544-4fbb-be7f-7081ff14ed3 6451849404952750 366586364589156 3778904.0 54667 176540398942621 2017-02-08 22:38:07 GENUINE
00046135-56ad-4ld4-bde8-3c51010c3547 343330859464924 051518771748227 5014013.0 81089 860541878277349 2016-02-04 07:01:07 GENUINE
004d284-clbc-496c-9ecc-d70bf6b78390 3492543308765970 175914389419795 5887044.0 50544 725031279093414 2017-05-11 09:07:59 GENUINE
00063233-8ff7-40bc-a74a-47826c594d8a 377201318164757 924475891017022 6881323.0 95975 749826149212397 2018-07-28 06:40:31 GENUINE
0008d2a6-c996-4ee2-83f7-014851a7a6fb 4907253800863053 723132659200637 5488105.0 23117 617446577958996 2017-05-08 23:11:54 GENUINE
0008dedef-1f7c-432b-9ed0-d2ac1446284f 6011920623387365 421308864367180 6482829.0 99114 28654301191234 2018-12-02 10:22:18 GENUINE
0009d330-bd31-456b-bf6c-7ab5b1dfb4b2 4546430394425179 459011730675059 2788646.0 52720 150020495648797 2017-06-07 07:11:35 GENUINE
000a3754-f342-42cb-9976-2c9162d03402 6227994101600953 940184607999133 5002389.0 21555 972412969980550 2017-01-04 18:11:33 GENUINE
000b1f21-c498-4fd0-9eb8-bb9ed6f9ade2 344583480345238 667772346348128 1982685.0 52254 6083912611953205 2017-04-06 16:58:46 GENUINE
000bdfa1-3f51-4d8b-a9ee-2039ea5db1f4 345949260434768 317939112617073 8462321.0 12124 158131710971474 2017-04-29 14:28:54 GENUINE
000ebe74-8279-4de1-800a-03ce39620262 4250028739827574 394067902581867 5430837.0 33820 491369407491686 2017-09-07 09:29:30 GENUINE
00118671-5ae9-42ca-bfdb-a5d1e3d5db6 4482511371042362 66599708492282 6761452.0 49620 790484082532571 2017-10-26 13:15:13 GENUINE
00121669-4f35-4967-95e2-48616a04d472 6451387765835527 188977299867156 3086412.0 99761 028580296523857 2017-12-05 01:16:36 GENUINE
00124d6c-0526-44c8-8f27-98e5ac62fc5e 5381466904123396 550262364037847 1136800.0 58269 715722111477598 2017-02-09 17:56:48 GENUINE
001312a6-c93c-45cc-873d-0fc3cc6f3548 4750696980073601 677295725087734 1187814.0 51246 767313351253016 2017-07-20 07:29:06 GENUINE
00144543-b32c-4d5c-bce2-8a97e354eac8 347274385563543 359453922401698 2757804.0 97133 183518215701679 2016-03-15 23:07:53 GENUINE
00184dd8-2c9d-4304-ae89-76c71924fadb 5252551880815473 356544959463542 5199495.0 83101 025098568987945 2017-09-14 04:45:44 GENUINE
001854db-e5f2-4dc9-b9be-c7463a45008e 6011980249342475 278094439476476 4436433.0 30144 816262392083518 2017-12-05 04:47:17 GENUINE
0019695d-8a44-447c-94c6-f2c622713d54 342679959520647 328980632419725 4342881.0 59867 64157458024598 2017-09-01 02:30:36 GENUINE
Time taken: 0.668 seconds, Fetched: 20 row(s)
hive>
```

### Below screenshot shows count of records in card\_transactions\_hbase table

```
hive> select count(*) from card_transactions_hbase;
Query ID = ec2-user_20190601114545_57ab7485-a9b2-474c-88de-b07a6e378e03
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1559386504130_0003, Tracking URL = http://ip-172-31-91-95.ec2.internal:8088/proxy/application_1559386504130_0003/
Kill Command = /opt/cloudera/parcels/CDH-5.15.0-1.cdh5.15.0.p0.21/lib/hadoop/bin/hadoop job -kill job_1559386504130_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2019-06-01 11:46:00,586 Stage-1 map = 0%,  reduce = 0%
2019-06-01 11:46:07,924 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 5.99 sec
2019-06-01 11:46:14,167 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 9.11 sec
MapReduce Total cumulative CPU time: 9 seconds 110 msec
Ended Job = job_1559386504130_0003
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1  Cumulative CPU: 9.11 sec  HDFS Read: 9416 HDFS Write: 6 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 110 msec
OK
59367
Time taken: 31.512 seconds, Fetched: 1 row(s)
hive>
```

### Below screenshot shows some records from lookup\_data\_hbase table

```

hive> select * from lookup_data_hbase limit 20;
OK
+-----+-----+-----+-----+
|      |      |      |      |
| 340028465709212 1.6331555548882348E7 | 233  | 35775 | 2018-11-22 11:08:25 |
| 340054675199675 1.4156079786189131E7 | 631  | 54499 | 2018-08-30 00:24:04 |
| 340082915339645 1.5285685330791473E7 | 407  | 75772 | 2018-08-17 18:35:42 |
| 340134186926007 1.5239767522438556E7 | 614  | 65055 | 2018-10-30 11:24:46 |
| 340265728490548 1.608491671255562E7 | 202  | 72411 | 2018-05-28 12:51:31 |
| 340268219434811 1.2507323937605347E7 | 415  | 73041 | 2018-07-29 11:49:54 |
| 340379737226464 1.4198310998368107E7 | 229  | 18054 | 2018-05-30 15:41:44 |
| 340383645652108 1.4091750460468251E7 | 645  | 98027 | 2018-10-29 14:28:48 |
| 340803866934451 1.0843341196185412E7 | 502  | 64642 | 2018-04-24 09:29:27 |
| 340889618969736 1.3217942365515321E7 | 330  | 28009 | 2018-06-29 17:07:40 |
| 340924125838453 1.1712496575647112E7 | 644  | 45880 | 2018-08-31 09:17:44 |
| 341005627432127 1.2843150457220368E7 | 629  | 70710 | 2018-09-27 13:56:56 |
| 341029651579925 1.384435212566918E7 | 682  | 84750 | 2018-11-28 01:42:28 |
| 341311317050937 1.2104977216410311E7 | 440  | 31066 | 2018-02-26 01:54:36 |
| 341344252914274 1.0770448614621945E7 | 643  | 19432 | 2018-05-27 22:39:16 |
| 341363858179050 1.2514602030066479E7 | 559  | 15685 | 2018-09-28 04:09:12 |
| 341519629171378 1.529211229141615E7 | 408  | 72377 | 2018-10-31 10:35:18 |
| 341641153427489 1.3421215715089727E7 | 475  | 64682 | 2018-09-25 20:05:00 |
| 341719092861087 1.1521372646859605E7 | 442  | 36740 | 2018-11-18 21:54:51 |
| 341722035429601 1.2890883282580867E7 | 552  | 15488 | 2018-06-23 16:05:37 |
+-----+-----+-----+-----+
Time taken: 0.096 seconds, Fetched: 20 row(s)
hive> 

```

**Below screenshot shows count of records in lookup\_data\_hbase table**

```

hive> select count(*) from lookup_data_hbase;
Query ID = ec2-user_20190601114848_le534600-e31e-4b0f-a8d2-2354bc4cd7e8
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1559386504130_0004, Tracking URL = http://ip-172-31-91-95.ec2.internal:8088/proxy/application_1559386504130_0004/
Kill Command = /opt/cloudera/parcels/CDH-5.15.0-1.cdh5.15.0.p0.21/lib/hadoop/bin/hadoop job -kill job_1559386504130_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2019-06-01 11:48:35,683 Stage-1 map = 0%,  reduce = 0%
2019-06-01 11:48:43,038 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 4.09 sec
2019-06-01 11:48:49,269 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 6.95 sec
MapReduce Total cumulative CPU time: 6 seconds 950 msec
Ended Job = job_1559386504130_0004
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1   Cumulative CPU: 6.95 sec   HDFS Read: 8799 HDFS Write: 4 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 950 msec
OK
999
Time taken: 29.962 seconds, Fetched: 1 row(s)
hive> 

```

**Below screenshot shows count of records from card\_transactions\_hbase where status is GENUINE**

```

hive> select count(*) from card_transactions_hbase where status = 'GENUINE';
Query ID = ec2-user_20190601115151_fc8a73d1-b604-43fc-8ece-1f7a20c4d441
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1559386504130_0005, Tracking URL = http://ip-172-31-91-95.ec2.internal:8088/proxy/application_1559386504130_0005/
Kill Command = /opt/cloudera/parcels/CDH-5.15.0-1.cdh5.15.0.p0.21/lib/hadoop/bin/hadoop job -kill job_1559386504130_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2019-06-01 11:51:19,562 Stage-1 map = 0%, reduce = 0%
2019-06-01 11:51:27,845 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 7.2 sec
2019-06-01 11:51:35,176 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 10.26 sec
MapReduce Total cumulative CPU time: 10 seconds 260 msec
Ended Job = job_1559386504130_0005
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 10.26 sec HDFS Read: 10467 HDFS Write: 6 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 260 msec
OK
59150
Time taken: 32.119 seconds, Fetched: 1 row(s)
hive> 

```

**Below screenshot shows count of records from card\_transactions\_hbase table where status is FRAUD**

```

hive> select count(*) from card_transactions_hbase where status = 'FRAUD';
Query ID = ec2-user_20190601115252_909d487e-d445-4328-8bea-6d50cbc2e33c
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1559386504130_0006, Tracking URL = http://ip-172-31-91-95.ec2.internal:8088/proxy/application_1559386504130_0006/
Kill Command = /opt/cloudera/parcels/CDH-5.15.0-1.cdh5.15.0.p0.21/lib/hadoop/bin/hadoop job -kill job_1559386504130_0006
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2019-06-01 11:52:31,317 Stage-1 map = 0%, reduce = 0%
2019-06-01 11:52:39,627 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 6.71 sec
2019-06-01 11:52:45,818 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 9.8 sec
MapReduce Total cumulative CPU time: 9 seconds 800 msec
Ended Job = job_1559386504130_0006
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 9.8 sec HDFS Read: 10514 HDFS Write: 4 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 800 msec
OK
217
Time taken: 33.419 seconds, Fetched: 1 row(s)
hive> 

```

**Showing processing of transactions for a card id – 6599900931314251, This is the same card id for which I had shown screenshots in mid submission as well. In continuation, I picked up the same card id.**

**Below are the transaction details for this card id received from Kafka topic captured in Output.txt. There are 7 transactions to be processed for this card id.**

```

{"card_id":6599900931314251,"member_id":928036864799687,"amount":1447809,"postcode":24848,"pos_id":302116371961141,"transaction_dt":"02-10-2018 12:49:01"}
{"card_id":6599900931314251,"member_id":928036864799687,"amount":9723817,"postcode":72043,"pos_id":239218493122711,"transaction_dt":"08-10-2018 10:06:47"}
{"card_id":6599900931314251,"member_id":928036864799687,"amount":6902536,"postcode":99739,"pos_id":97601123545965,"transaction_dt":"12-02-2018 12:56:07"}
{"card_id":6599900931314251,"member_id":928036864799687,"amount":5471142,"postcode":61237,"pos_id":430322863074336,"transaction_dt":"09-03-2018 20:30:08"}
{"card_id":6599900931314251,"member_id":928036864799687,"amount":4361135,"postcode":97026,"pos_id":346155771351832,"transaction_dt":"12-10-2018 03:49:30"}
 {"card_id":6599900931314251,"member_id":928036864799687,"amount":4876716,"postcode":49910,"pos_id":683867473927303,"transaction_dt":"21-11-2018 05:59:02"}
 {"card_id":6599900931314251,"member_id":928036864799687,"amount":2179073,"postcode":34270,"pos_id":229855333287383,"transaction_dt":"27-06-2018 14:17:28"} 

```

**Below are the screenshots of processing of transactions for this card id captured in Output.txt**

**Transaction 1:**

```
=====
New Transaction Processing Start Time : 2019-06-01 11:15:42
===== Classify Current Transaction into GENUINE/FRAUD =====
Card ID of current transaction is : 6599900931314251
UCL (Upper Control Limit) for Card ID of current transaction is : 1.2121408572464656E7
Current Transaction Amount is : 1447809.0
Credit score for Card ID of current transaction is : 297
Last Post Code for Card ID of current transaction is : 97423
Current Post Code for Card ID of current transaction is : 24848
Last Transaction Date for Card ID of current transaction is : 2018-01-31 11:25:16
Current Transaction Date for Card ID of current transaction is : 2018-10-02 12:49:01
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 3644.2925949321507
Absolute Time Difference between current and last transaction (in seconds) is : 21086625
Speed (Distance / Time) : 1.7282484015019713E-4 km/second
Current Card Transaction has passed Credit Score Validation rule : Score : 297 >= 200
Current Card Transaction has passed UCL (Upper Control Limit) Validation rule : Amount : 1447809.0 <= 1.2121408572464656E7
Current Card Transaction has passed Zip Code Distance Validation rule : Speed (km/sec) : 1.7282484015019713E-4 <= 0.25
Current Card Transaction has passed all 3 Validation rules and so status is : GENUINE
===== Update NoSQL Database for Current Transaction =====
Card ID of current transaction is : 6599900931314251
Current Transaction Details are populated in Card Transactions HBase Table : card_transactions_hive
Current Card Transaction status is GENUINE so updating lookup table
Postcode and Transaction Date updated for Card ID of current transaction in lookup HBase table : lookup_data_hive
=====
Current Transaction Processing End Time : 2019-06-01 11:15:42
=====
```

**Transaction 2:**

```
=====
New Transaction Processing Start Time : 2019-06-01 11:15:42
===== Classify Current Transaction into GENUINE/FRAUD =====
Card ID of current transaction is : 6599900931314251
UCL (Upper Control Limit) for Card ID of current transaction is : 1.2121408572464656E7
Current Transaction Amount is : 9723817.0
Credit score for Card ID of current transaction is : 297
Last Post Code for Card ID of current transaction is : 24848
Current Post Code for Card ID of current transaction is : 72043
Last Transaction Date for Card ID of current transaction is : 2018-10-02 12:49:01
Current Transaction Date for Card ID of current transaction is : 2018-10-08 10:06:47
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 892.4283020933965
Absolute Time Difference between current and last transaction (in seconds) is : 508666
Speed (Distance / Time) : 0.0017544485027373492 km/second
Current Card Transaction has passed Credit Score Validation rule : Score : 297 >= 200
Current Card Transaction has passed UCL (Upper Control Limit) Validation rule : Amount : 9723817.0 <= 1.2121408572464656E7
Current Card Transaction has passed Zip Code Distance Validation rule : Speed (km/sec) : 0.0017544485027373492 <= 0.25
Current Card Transaction has passed all 3 Validation rules and so status is : GENUINE
===== Update NoSQL Database for Current Transaction =====
Card ID of current transaction is : 6599900931314251
Current Transaction Details are populated in Card Transactions HBase Table : card_transactions_hive
Current Card Transaction status is GENUINE so updating lookup table
Postcode and Transaction Date updated for Card ID of current transaction in lookup HBase table : lookup_data_hive
=====
Current Transaction Processing End Time : 2019-06-01 11:15:42
=====
```

**Transaction 3:**

```
=====
New Transaction Processing Start Time : 2019-06-01 11:15:42
===== Classify Current Transaction into GENUINE/FRAUD =====

Card ID of current transaction is : 6599900931314251
UCL (Upper Control Limit) for Card ID of current transaction is : 1.2121408572464656E7
Current Transaction Amount is : 6902536.0
Credit score for Card ID of current transaction is : 297
Last Post Code for Card ID of current transaction is : 72043
Current Post Code for Card ID of current transaction is : 99739
Last Transaction Date for Card ID of current transaction is : 2018-10-08 10:06:47
Current Transaction Date for Card ID of current transaction is : 2018-02-12 12:56:07
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 5598.51251087252
Absolute Time Difference between current and last transaction (in seconds) is : 20553040
Speed (Distance / Time) : 2.723934031594606E-4 km/second
Current Card Transaction has passed Credit Score Validation rule : Score : 297 >= 200
Current Card Transaction has passed UCL (Upper Control Limit) Validation rule : Amount : 6902536.0 <= 1.2121408572464656E7
Current Card Transaction has passed Zip Code Distance Validation rule : Speed (km/sec) : 2.723934031594606E-4 <= 0.25
Current Card Transaction has passed all 3 Validation rules and so status is : GENUINE
===== Update NoSQL Database for Current Transaction =====

Card ID of current transaction is : 6599900931314251
Current Transaction Details are populated in Card Transactions HBase Table : card_transactions_hive
Current Card Transaction status is GENUINE so updating lookup table
Postcode and Transaction Date updated for Card ID of current transaction in lookup HBase table : lookup_data_hive
=====

Current Transaction Processing End Time : 2019-06-01 11:15:42
=====
```

**Transaction 4:**

```
=====
New Transaction Processing Start Time : 2019-06-01 11:15:42
===== Classify Current Transaction into GENUINE/FRAUD =====
Card ID of current transaction is : 6599900931314251
UCL (Upper Control Limit) for Card ID of current transaction is : 1.2121408572464656E7
Current Transaction Amount is : 5471142.0
Credit score for Card ID of current transaction is : 297
Last Post Code for Card ID of current transaction is : 99739
Current Post Code for Card ID of current transaction is : 61237
Last Transaction Date for Card ID of current transaction is : 2018-02-12 12:56:07
Current Transaction Date for Card ID of current transaction is : 2018-03-09 20:30:08
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 5082.797030871786
Absolute Time Difference between current and last transaction (in seconds) is : 2187241
Speed (Distance / Time) : 0.002323839499566708 km/second
Current Card Transaction has passed Credit Score Validation rule : Score : 297 >= 200
Current Card Transaction has passed UCL (Upper Control Limit) Validation rule : Amount : 5471142.0 <= 1.2121408572464656E7
Current Card Transaction has passed Zip Code Distance Validation rule : Speed (km/sec) : 0.002323839499566708 <= 0.25
Current Card Transaction has passed all 3 Validation rules and so status is : GENUINE
===== Update NoSQL Database for Current Transaction =====
Card ID of current transaction is : 6599900931314251
Current Transaction Details are populated in Card Transactions HBase Table : card_transactions_hive
Current Card Transaction status is GENUINE so updating lookup table
Postcode and Transaction Date updated for Card ID of current transaction in lookup HBase table : lookup_data_hive
=====

Current Transaction Processing End Time : 2019-06-01 11:15:42
=====
```

**Transaction 5:**

```
=====
New Transaction Processing Start Time : 2019-06-01 11:15:42
===== Classify Current Transaction into GENUINE/FRAUD =====
Card ID of current transaction is : 6599900931314251
UCL (Upper Control Limit) for Card ID of current transaction is : 1.2121408572464656E7
Current Transaction Amount is : 4361135.0
Credit score for Card ID of current transaction is : 297
Last Post Code for Card ID of current transaction is : 61237
Current Post Code for Card ID of current transaction is : 97026
Last Transaction Date for Card ID of current transaction is : 2018-03-09 20:30:08
Current Transaction Date for Card ID of current transaction is : 2018-10-12 03:49:30
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 2616.258781260241
Absolute Time Difference between current and last transaction (in seconds) is : 18688762
Speed (Distance / Time) : 1.3999101605875451E-4 km/second
Current Card Transaction has passed Credit Score Validation rule : Score : 297 >= 200
Current Card Transaction has passed UCL (Upper Control Limit) Validation rule : Amount : 4361135.0 <= 1.2121408572464656E7
Current Card Transaction has passed Zip Code Distance Validation rule : Speed (km/sec) : 1.3999101605875451E-4 <= 0.25
Current Card Transaction has passed all 3 Validation rules and so status is : GENUINE
===== Update NoSQL Database for Current Transaction =====
Card ID of current transaction is : 6599900931314251
Current Transaction Details are populated in Card Transactions HBase Table : card_transactions_hive
Current Card Transaction status is GENUINE so updating lookup table
Postcode and Transaction Date updated for Card ID of current transaction in lookup HBase table : lookup_data_hive
=====

Current Transaction Processing End Time : 2019-06-01 11:15:42
=====
```

**Transaction 6:**

```
=====
New Transaction Processing Start Time : 2019-06-01 11:15:42
===== Classify Current Transaction into GENUINE/FRAUD =====
Card ID of current transaction is : 6599900931314251
UCL (Upper Control Limit) for Card ID of current transaction is : 1.2121408572464656E7
Current Transaction Amount is : 4876716.0
Credit score for Card ID of current transaction is : 297
Last Post Code for Card ID of current transaction is : 97026
Current Post Code for Card ID of current transaction is : 49910
Last Transaction Date for Card ID of current transaction is : 2018-10-12 03:49:30
Current Transaction Date for Card ID of current transaction is : 2018-11-21 05:59:02
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 2567.88379389069
Absolute Time Difference between current and last transaction (in seconds) is : 3463772
Speed (Distance / Time) : 7.41354740984883E-4 km/second
Current Card Transaction has passed Credit Score Validation rule : Score : 297 >= 200
Current Card Transaction has passed UCL (Upper Control Limit) Validation rule : Amount : 4876716.0 <= 1.2121408572464656E7
Current Card Transaction has passed Zip Code Distance Validation rule : Speed (km/sec) : 7.41354740984883E-4 <= 0.25
Current Card Transaction has passed all 3 Validation rules and so status is : GENUINE
===== Update NoSQL Database for Current Transaction =====
Card ID of current transaction is : 6599900931314251
Current Transaction Details are populated in Card Transactions HBase Table : card_transactions_hive
Current Card Transaction status is GENUINE so updating lookup table
Postcode and Transaction Date updated for Card ID of current transaction in lookup HBase table : lookup_data_hive
=====

Current Transaction Processing End Time : 2019-06-01 11:15:43
=====
```

## Transaction 7:

```
=====
New Transaction Processing Start Time : 2019-06-01 11:15:43
===== Classify Current Transaction into GENUINE/FRAUD =====
Card ID of current transaction is : 6599900931314251
UCL (Upper Control Limit) for Card ID of current transaction is : 1.2121408572464656E7
Current Transaction Amount is : 2179073.0
Credit score for Card ID of current transaction is : 297
Last Post Code for Card ID of current transaction is : 49910
Current Post Code for Card ID of current transaction is : 34270
Last Transaction Date for Card ID of current transaction is : 2018-11-21 05:59:02
Current Transaction Date for Card ID of current transaction is : 2018-06-27 14:17:28
Distance (returned by DistanceUtility) between current and last postcode (in km) is : 2222.340348620956
Absolute Time Difference between current and last transaction (in seconds) is : 12670894
Speed (Distance / Time) : 1.7538938835893946E-4 km/second
Current Card Transaction has passed Credit Score Validation rule : Score : 297 >= 200
Current Card Transaction has passed UCL (Upper Control Limit) Validation rule : Amount : 2179073.0 <= 1.2121408572464656E7
Current Card Transaction has passed Zip Code Distance Validation rule : Speed (km/sec) : 1.7538938835893946E-4 <= 0.25
Current Card Transaction has passed all 3 Validation rules and so status is : GENUINE
===== Update NoSQL Database for Current Transaction =====
Card ID of current transaction is : 6599900931314251
Current Transaction Details are populated in Card Transactions HBase Table : card_transactions_hive
Current Card Transaction status is GENUINE so updating lookup table
Postcode and Transaction Date updated for Card ID of current transaction in lookup HBase table : lookup_data_hive
=====
Current Transaction Processing End Time : 2019-06-01 11:15:43
=====
```

Below are screenshots from HBase for this card id:

First screenshot shows postcode and transaction\_dt for same card id in lookup table

You will notice 10 VERSIONS for postcode and transaction\_dt. This column family is set to store up to 10 VERSIONS so all 10 VERSIONS are displayed. 3 versions came at the time of mid submission because I had populated lookup table 3 times. First, when I populated lookup table for first time in task 3. Second, when I populated lookup table for second time in task 4 while running oozie job without coordinator. Third, when I populated lookup table for third time in task 4 again while running oozie job with coordinator this time. Rest 7 versions came from processing of transactions for this card id. There were 7 transactions fetched for Kafka topic for this card and all 7 were GENUINE as shown in screenshots above already.

```

hbase(main):003:0> get 'lookup_data_hive', '6599900931314251', {COLUMN => ['lookup_transaction_family:postcode', 'lookup_transaction_family:transaction_dt'], VERSIONS=>10}
COLUMN
lookup_transaction_family:postcode
CELL
timestamp=1559387743032, value=34270
timestamp=1559387743012, value=49910
timestamp=1559387742983, value=97026
timestamp=1559387742962, value=61237
timestamp=1559387742941, value=99739
timestamp=1559387742920, value=72043
timestamp=1559387742899, value=24948
timestamp=1558807804067, value=97423
timestamp=1558774131995, value=97423
timestamp=1558772840191, value=97423
timestamp=1559387743032, value=2018-06-27 14:17:28
timestamp=1559387743012, value=2018-11-21 05:59:02
timestamp=1559387742983, value=2018-10-12 03:49:30
timestamp=1559387742962, value=2018-03-09 20:30:08
timestamp=1559387742941, value=2018-02-12 12:56:07
timestamp=1559387742920, value=2018-10-08 10:06:47
timestamp=1559387742899, value=2018-10-02 12:49:01
timestamp=1558807804067, value=2018-01-31 11:25:16
timestamp=1558774131995, value=2018-01-31 11:25:16
timestamp=1558772840191, value=2018-01-31 11:25:16
20 row(s) in 0.0190 seconds
hbase(main):004:0>

```

**Below screenshot shows ucl and score for same card id in lookup table. This column family is set to store up to 1 VERSION only so only 1 VERSION is displayed.**

```

hbase(main):004:0> get 'lookup_data_hive', '6599900931314251', {COLUMN => ['lookup_card_family:ucl', 'lookup_card_family:score'], VERSIONS=>10}
COLUMN
lookup_card_family:score
CELL
timestamp=1558807804067, value=297
lookup_card_family:ucl
CELL
timestamp=1558807804067, value=1.2121408572464656E7
2 row(s) in 0.0080 seconds
hbase(main):005:0>

```

**Below screenshot shows latest version of all columns for this card id in lookup table**

```

hbase(main):005:0> get 'lookup_data_hive', '6599900931314251'
COLUMN
lookup_card_family:score
CELL
timestamp=1558807804067, value=297
lookup_card_family:ucl
CELL
timestamp=1558807804067, value=1.2121408572464656E7
lookup_transaction_family:postcode
CELL
timestamp=1559387743032, value=34270
lookup_transaction_family:transaction_dt
CELL
timestamp=1559387743032, value=2018-06-27 14:17:28
4 row(s) in 0.0070 seconds
hbase(main):006:0>

```

===== Output – END =====

===== Final Submission – END =====